

CONVERSATIONS WITH CLAUDE

VOLUME ONE

THE INFINITE

by Suraaj Chandra Jha

a free book about talking to machines

THIS BOOK IS FREE

Take it. Give it away.

This book is free. It always will be. If someone charged you for it, you have been scammed, and honestly, ask them for their marketing secrets. Send this file to anyone: mail it, post it, print it. Just keep it whole and keep my name on it. The lawyers call that CC BY-NC-ND 4.0. I call it manners.

Full screen. Press next.

Read it full screen. Press next. One idea per page, in order, cover to cover, in about an evening. Or skip around: every page stands alone. There is no wrong way, but full screen with a beverage is the intended experience.

Not the manufacturer.

Claude is made by Anthropic. This book is not: not affiliated, not endorsed, not reviewed. I am a hacker who spends an unreasonable amount of time talking to their machine and taking notes. Software changes faster than books. Where a detail has drifted since printing, trust the live documentation and forgive the author.

Ten doors.

01 READ THIS FIRST

02 WHAT IS THIS THING

03 FIRST CONTACT

04 IT IS NOT GOOGLE

05 HOW TO ASK

06 THE FOLLOW-UP

07 WHAT IT IS GREAT AT

08 WHAT IT IS BAD AT

09 TRUST, BUT VERIFY

10 RECIPES

READ THIS FIRST

Nobody reads manuals. This is not a manual.

This is a conversation.

Every page of this book holds exactly one idea, set large enough to read from across the room. You turn the page, you get the next idea. There is no chapter you must read first, no appendix doing secret load-bearing work. Read it straight through in an evening, or wander. Manuals are written for the machine, in the machine's order. This book is written for you, in the order curiosity moves.

We assume you know nothing.

That is a compliment, not an insult. Explanations written for experts get to cheat: a nod, an acronym, a knowing shrug. Explanations written for beginners have to actually explain, so beginners get the good ones. Zero is a fine place to start. Every expert did.

Meet Claude.

Claude is an AI assistant made by a company called Anthropic. AI is short for artificial intelligence: software doing the kind of work that used to need a mind. You type into a box, in ordinary language, and it types back: answers, drafts, explanations, ideas. It lives at claude.ai in your browser and in apps on your phone and computer. That is everything you need to know to begin, and half the software industry would rather you believed otherwise. What an AI assistant is, and how a machine came to hold up its end of a conversation, is the next chapter's business. For now: a box, you, and something on the other side that answers.

I am Suraaj. I am a hacker.

Before you close the book: a hacker is a person who delights in understanding a system and making it do wonderful things it was not obviously built to do. The word meant exactly that for decades before the newspapers borrowed it for people who break into banks. Those people already had a word, criminals. I have spent my life poking at machines to see what they will do, and I once taught my aunt to use email, one Sunday at a time. This machine is the most interesting one I have ever poked. These three volumes are the trip report.

YOU:

whats a hacker? my uncle
says its ppl who steal
passwords

CLAUDE:

Originally, someone who
explores systems for the joy
of it. Your uncle means
criminals.

Try things as you read.

This book has a single rule. Open Claude in another window, right now, and each time a page hands you an idea, walk it over and try it. The first visit to `claude.ai` asks you to create an account: free, and about a minute of your life. Reading about conversation is like reading about swimming: accurate, and dry. The book is the map. The chat box is the territory. You cannot break anything by typing, a later page makes that case properly, and the entire difference between reading about this machine and being good with it is thirty seconds of trying, repeated.

YOU:

ok the book told me to try something. um. why is the sky blue but sunsets are red

CLAUDE:

Same air, longer path. The blue scatters away before the light reaches you.

Why this volume is called THE INFINITE.

Not because the machine is infinite. It is a made thing, with limits this book will happily point at. The infinite part is not in the machine at all. It sits somewhere much closer to your chair, and naming it now would spoil it: the last page of this volume points straight at it, and by then you will have earned the view. What you do not know yet is what is answering when you type. Turn the page. It is stranger and simpler than you think.

WHAT IS THIS THING

Every sufficiently advanced autocomplete is indistinguishable from magic.

It finishes text for a living.

Given some text, it guesses what comes next. One word, then the next, then the next. Done at absurd scale, with guesses this good, the trick produces essays, working code, and wedding toasts. Everything else in this book is a footnote to this page.

Better read than anyone you know.

To get good at guessing, it first did the reading: a huge slice of everything humans have written down. Books, articles, conversations, code, great swaths of the public web. The process is called training, and one detail matters more than the rest: none of that text sits inside it like files in a folder. It was absorbed, the way a lifetime of reading is absorbed by you. You do not keep a copy of every book you have finished. You keep the habits they left behind. So does it.

YOU:

finish this sentence: it was
a dark and stormy

CLAUDE:

night. Give me a harder one.

Nothing in there is looked up.

A database stores facts and hands them back exactly as filed. This machine does neither. When you ask a question, the answer is assembled on the spot, as you would retell a story you know well: mostly right, occasionally embroidered with a detail that was never there, and delivered in the same confident voice either way. Plant a flag on this page. It explains the machine's strangest failures, and Volume Two returns to climb it properly. For now, remember that every answer is a performance, not a printout.

YOU:

whats the first line of moby dick

CLAUDE:

Call me Ishmael. For line two hundred, I would be improvising.

Billions of tiny dials.

Inside the machine are billions of numbers called parameters. Picture a wall of tiny dials stretching past the horizon. Training works like this: the machine guesses the next word of a real sentence, peeks at the actual next word, and every dial gets nudged a hair in the direction that would have improved the guess. Repeat, a stupendous number of times. No engineer ever set a single dial by hand, and nobody can point to the dial that holds Paris or the past tense. The knowledge is not in any one dial. It lives in the pattern of all of them.

★ The people who built it cannot read the dials either. It bothers them too.

Why it is called a model.

The technical name for this thing is a language model, and the second word is doing honest work. A globe is a model of the Earth: dramatically smaller, wrong about plenty of details, and still the best way to see how everything fits together. A language model strikes the same bargain with everything we have written: compressed, and useful out of all proportion to its faults. Nobody sues the globe for leaving out their street. The skill is knowing what to ask of it, and what to go check somewhere else.

So is it thinking?

Nobody knows. The people who build these machines argue about it in public. Some say prediction at this scale adds up to a kind of thinking. Others say it is an imitation so good the difference stops mattering, or never existed. The field has not settled the question, so this book will not pretend to. You do not need the answer. A hammer drives nails with or without a theory of carpentry. Use the machine, watch what it does well and badly. The philosophers will keep their jobs either way.

YOU:

be honest, are you actually thinking right now

CLAUDE:

Something is happening.
Naming it is above both our
pay grades.

Meet the token.

One term of art before we go, because you will meet it everywhere: a token is a chunk of text, usually a short word or a piece of a longer one, and tokens are what the machine reads and writes. Not letters. Chunks. That one fact explains some odd behavior you will notice later, like why counting the letters in a word can trip it up. The full story gets its own chapter in Volume Two. Around here we defer things. We do not hand-wave them.

Claude, by Anthropic, tuned for conversation.

Claude is what this book talks to: a family of these models built by a company called Anthropic, then tuned, on purpose, to be helpful and honest in conversation rather than to merely finish sentences. Other assistants exist, built by other companies on the same core trick, and nearly everything in these three volumes transfers to them. I write about Claude because Claude is where I live: its habits are the ones I know. That is the anatomy lesson, done. Next chapter you stop reading about the machine and say something to it instead.

FIRST CONTACT

Say hi. It doesn't bite. It doesn't anything.

The blank box is the whole interface.

Open Claude and you get a box. No menus, no settings, no tutorial demanding your attention first. Decades of software taught you that power hides behind complexity: the more buttons, the more it must do. This is the opposite machine. The empty box is not a search bar waiting for keywords, and not a command line, the old screen that does nothing until a programmer types the spell exactly right. It is an opening shaped like a person. Say what you want the way you would say it out loud, and it goes to work. The hardest part of starting is believing it is really that plain.

Typos do not count against you.

You do not need to write well to be understood here. Misspell things. Abandon grammar halfway through. Ask in whatever language you think in, or in two of them tangled together mid-sentence. Claude has read enormous amounts of real human writing, and real human writing is a mess, so mess is a dialect it speaks fluently. It answers the question you meant, not the one your typos wrote. Nobody is grading you. The red squiggle that has followed you around since school does not live in this box.

YOU:

how do i make my wifi more
faster it keep stoping when
i go kichen

CLAUDE:

Sounds like range. Move the
router higher and more
central, then we talk
extenders.

There is no wrong button.

Every program before this one taught you caution: one wrong click deletes the file, sends the email early. A conversation has no wrong click. The worst case here is a dull answer, and the fix is asking again, differently. You cannot break it. Poke it. Curiosity is the intended input.

It speaks every register you can name.

Register is the word for how a thing is said: boardroom formal, kitchen-table casual, legal, poetic, playground. Humans spend years learning to switch registers, and most of us manage three. Claude switches on request, instantly, on any subject. The same fact can arrive as a memo, a bedtime story, or a sea shanty, and the fact survives the costume change. Matching the register to the reader is half of communication, and the dial is now in your hands. Ask for the voice you need.

YOU:

why does bread rise. formal first, then how my gran would say it

CLAUDE:

Yeast metabolizes sugars and releases carbon dioxide. Or: the yeast burps, love, and the dough puffs right up.

“Explain it simpler” is a legal move.

Somewhere in your past, a teacher explained a thing twice, and the third ask felt too expensive. Every learner knows that arithmetic: each repeat costs goodwill, so you nod and stay lost. That economy does not exist here. Say “simpler” and you get simpler. Say “simpler than that” and it goes down another level, through analogies and examples, as many floors as the building has, and the tenth request is received exactly like the first. For explanation, infinite patience is the entire product. No human explainer can offer it, because humans run out.

Every answer is a first draft.

A search result is final: take it or scroll on. A Claude answer is an opening offer, the first position in a negotiation you are always allowed to continue. Too long, say shorter. Wrong angle, say which angle. Beginners read the first reply as a verdict and either accept it or give up, and both moves waste the machine. Nothing in the conversation is finished until you say it is. The reply on your screen is not the answer. It is the start of one.

Say “no, more like this”.

Rejecting a colleague’s third draft takes diplomacy, timing, and possibly a coffee. Rejecting Claude’s third draft takes four words. It does not sulk, defend the work, or remember the insult, because there is no insult: a correction is just more information about what you want, and it treats every one as a gift. Steer bluntly. “No, warmer.” “Shorter.” “More like the second one.” Each nudge lands in full, and nothing is held against you. The awkwardness you feel is yours alone, and it will pass.

YOU:

no, no, still no. write it
like an actual human who
likes me

CLAUDE:

Draft four: Hi Maya, quick
one about Friday. I need to
move our call.

Talking was the easy part.

That is everything you need to hold a conversation: say the thing, break any rule of grammar you like, steer bluntly, ask again forever. Notice what was missing: no syntax, no manual, no skill you did not already have from a lifetime of talking to people. But one old habit is waiting to ambush you. The last box you typed into this often belonged to a search engine, and your fingers will want to treat this one the same way. It is not the same kind of box at all. Next: why.

IT IS NOT GOOGLE

Stop googling at it.

Search finds. Claude makes.

A search engine finds: it hands you pages that already existed, written by strangers. Claude makes: the answer did not exist anywhere until you asked. It is composed on the spot, shaped to your exact question. Same box, different machinery. Confusing them is the classic beginner mistake.

Keywords are a habit to break.

Search taught you to speak in grunts. You wanted to know whether the ficus would survive winter, so you typed 'ficus cold tolerance'. The compression was for the machine's benefit: it could match words, not understand them. Claude understands sentences, so compressing only throws away information it could have used. Say the whole thought: where the plant sits, what the room is like, what you already tried. The more of your actual situation reaches it, the more the answer fits your actual situation.

YOU:

ficus dropping leaves why

CLAUDE:

A symptom with a dozen causes. Where does it live, and since when?

It went into the exam hall.

Claude learned what it knows by reading, and at some point the reading stopped. Picture a brilliant student who devoured the whole library, then walked into an exam hall and handed over its phone. Inside the hall it can reason beautifully about everything it read. But yesterday's election, this morning's weather, the game last night: none of that ever reached its desk. The stopping date is called the training cutoff. Every model has one, newer models have later ones, and the current date lives in Anthropic's docs, not here, because this page does not update itself.

But it can look things up.

The exam hall has a door. Ask Claude about something recent and it can step out: search the web, read the pages it finds, and come back with an answer that points at its sources. So there are two modes. Memory is it speaking from everything absorbed in training. Research is it reading live pages right now and citing them. The one skill to practice is noticing which mode you are in: links and sources mean research, a confident bare answer means memory.

YOU:

is the flower show still on
this weekend

CLAUDE:

Checked the organizer's page
just now: moved to next
Saturday, rain.

Current and exact: search. Thinking: converse.

When the answer must be current or exact, a train time, a price, a dosage, a deadline, use search, or ask Claude to search and show its sources. When what you want is understanding, a draft, or company for your thinking, converse. Most early disappointment comes from getting this backwards: asking a search engine to think, or asking Claude for train times from memory, then blaming the tools. Keep both tabs open.

The librarian and the friend.

A librarian hands you the exact book, open to the exact page, and vouches for nothing except where it was shelved. A well-read friend explains the whole subject over coffee, connects it to your situation, and now and then misremembers a detail with complete confidence. Neither replaces the other, and nobody sensible fires the librarian for not being warm, or the friend for not being an index. You have wanted both in your life since forever. Now you know which door to knock on.

YOU:

ok so which one are you

CLAUDE:

The friend. I can walk over to the librarian whenever you ask.

A conversation compounds. A search does not.

The deeper difference: a conversation accumulates. Ten searches are ten strangers, each starting from zero, each forgetting you on arrival. Ten exchanges with Claude are one thread: every answer becomes context for the next question, corrections stick, and by the tenth turn it is not answering trivia, it is helping with your project. Interest accrues on everything you have said so far. But compounding needs a deposit, and the deposit is how you ask. That is the next chapter, and it is the one that makes the rest of this book pay.

HOW TO ASK

It cannot read your mind. It has read nearly everything else.

Most bad answers are missing context.

Claude cannot see your desk, your inbox, your calendar, or the meeting you just walked out of. It knows your situation exactly as well as you have described it, which at the start of a conversation is not at all. Most disappointing answers are failures of briefing, not intelligence: you asked for directions without ever saying where you are standing. The fix costs one sentence. Before the question, say where you are and what you are trying to do. The answers sharpen immediately.

★ It is not politely ignoring your inbox. As far as it knows, you do not have one.

Describe the output, not just the topic.

“Write an email” is a topic, not a request. It leaves Claude guessing at length, tone, and stakes, so it guesses something safe and roomy: five polite paragraphs, hoping one of them is what you meant. “A three-line email to my landlord about the dead heater, friendly but firm” is a brief, and it comes back ready to send. Format is part of the ask. So are length and tone. None of this is advanced technique. It is the difference between telling the waiter you are hungry and ordering lunch.

YOU:

that came out five paragraphs. try again: 3 lines, to my landlord, friendly but firm

CLAUDE:

Hi Sam. The heater died Tuesday, could someone come this week? Thanks as always.

Audience is context too.

The same explanation has to be rebuilt for every listener. Your project is late: the version for your boss is two sentences and a new date, the version for your team names the blocker, the version for the customer apologizes without confessing. You already know this, it is half of why writing feels hard. Claude will do the rebuilding, but only if it knows who is listening. Add “for my boss” or “for a nervous new customer” to the ask. The answer changes shape.

YOU:

whats a 401k... actually
wait, explain it to my dad.
hes 70 and suspicious of
banks

CLAUDE:

Started with the part where
the money stays his, no
matter what the bank does.

One good example beats three paragraphs.

Trying to describe the style you want is hard. Showing it is easy. Paste something you like, an email that landed well, last month's report, a paragraph whose tone you envy, and say: like this, but for my situation. An example carries a hundred small decisions you would never think to spell out, sentence length, formality, where the warmth goes. Claude reads all of them at once. Show first, then tell. This one trick is worth the price of the whole volume. The price is zero. Still.

A brief's best word is no.

Tell it what it may not do. Under a hundred words. No jargon. Nothing that sounds like a press release. Given a wide-open ask, it produces the average of everything plausible, and averages are mush. Each constraint carves away mush. The sonnet's fourteen lines are why sonnets are good.

Give it one job at a time.

“Summarize this report, and make it funnier, and translate the summary into Spanish, and turn that into slides.” Each piece is easy. Braided together they tangle: the summary comes out half funny, the Spanish inherits the wrong draft, the slides inherit everything. While you are learning, give one job per ask. Look at the result, then hand over the next job. You lose nothing by going in steps, because the conversation remembers every step. Chaining jobs gracefully is a real skill, and it arrives on its own about one chapter from now.

Ask me questions before you answer.

Say it exactly like that. Instead of guessing at everything you forgot to mention, Claude interviews you. It asks about the audience, the deadline, the thing you were avoiding saying. You answer in lazy fragments, which is fine, fragments count. Then it writes. The result often does not need a second round, because the machine knows what it is missing better than you know what you left out. Six words, and the guessing stops.

YOU:

help me write my dads
retirement speech. ask me
questions before you answer

CLAUDE:

How long is the speech, who
is in the room, and what
story does your family
always tell about him?

Brief it like a colleague.

Nothing in this chapter was new to you. You brief people constantly: you tell the plumber which tap, you tell a colleague what the client meant, you give a friend the backstory before the question. The search box trained you to strip all of that away. Put it back. Talk to Claude the way you talk to a capable colleague on their first day: context, audience, an example, a few constraints. Then comes the part beginners miss. A good brief buys you a good first draft. What you say to the draft is where the real game starts.

THE FOLLOW-UP

The second prompt is the one that counts.

The first answer is the beginning.

You briefed it like a colleague. It handed you a draft. Beginners take that draft, say thanks in their head, and leave. The first reply is an opening offer, not the result. The real work lives in replies two through ten, where you push, choose, question, and shape.

It takes the note. Instantly.

Tell it what is wrong, plainly. Too formal. Wrong angle. You missed the point, I meant the other thing. The note lands instantly and completely, with no feelings involved, because there are none to involve. Half of what makes human feedback expensive is the diplomacy tax: softening, framing, timing. Here the tax is zero. Say the blunt thing. Blunt is not rude when nobody is on the receiving end.

YOU:

no offense but paragraph two
reads like a robot wrote it

CLAUDE:

None taken. None available,
actually. Rewritten.

Order three versions: safe, bold, weird.

Describing the perfect version is hard. Recognizing it is easy. So stop describing and start choosing: ask for three versions, one safe, one bold, one weird. The safe one tells you the floor. The weird one tells you where the walls are. The bold one is usually the winner, or close enough to point at and say, this one, but shorter. Editors and art directors have worked this way forever. You are allowed to work this way too, and the extra drafts cost you nothing but the asking.

YOU:

3 names for my sourdough
newsletter. safe, bold,
weird

CLAUDE:

The Weekly Loaf. Rise. The
Wild Yeast Chase.

Ask it what is weak.

When the draft is close but not right and you cannot name why, ask the author. What is weak about this draft? It critiques its own work surprisingly well: the vague second paragraph, the buried point, the ending that trails off. Then say, fix those things. Two short messages, one better draft. This sounds like asking the fox to audit the henhouse, but writing and judging are different acts, for machines as for people. Every writer you admire is also a ruthless reader of their own pages. The machine just skips the part where that hurts.

Interrogate the answer.

The conversation ends when you accept the answer. Question it instead, and a better one starts. Ask why. Ask for the counterargument. Ask what you are not asking. These moves work because the machine has read every side of nearly every argument and holds no stake in yours. Used this way, the conversation stops being a vending machine and becomes a thinking gym: you go there to get stronger at questions.

YOU:

ok what am i not asking that
i should be

CLAUDE:

Whether you want this job,
or just the offer.

Rough shape first. Then passes.

Nobody carves a statue in one stroke. A sculptor blocks out the rough shape, then makes passes: form, then surfaces, then detail. Work with the machine the same way. First message gets the rough shape. The next few are passes: fix the structure, then the tone, then the small stuff, one layer per message.

Expecting a finished piece from a single prompt, meaning one typed message, is not high standards, it is a misunderstanding of how making works, and it was one long before this machine showed up.

★ Michelangelo worked in passes too. His undo button was worse.

The conversation is the unit of work.

Somewhere in the last six pages the prompt stopped being the unit of work. The conversation is the unit of work: a first draft, a pushback, three options, a critique, a question, a pass. That is why this book is not called PROMPTS FOR CLAUDE. Prompts are single moves. Conversations are where the results live. The next question is where to point one, because some jobs suit this machine far better than others, and the next chapter is the tour.

WHAT IT IS GREAT AT

A partial list of superpowers, alphabetized by joy.

Text in, different text out.

Give it any text and describe the text you want instead. A rambling voice memo becomes a tidy summary. A long report becomes three sentences. English becomes Japanese, a wall of notes becomes a table. Nothing new has to be invented. The material is already there, it just changes shape. This is the workhorse superpower, the one you will use daily long after the wonder wears off. If a job can be phrased as “turn this into that”, it is already won.

A bad draft beats a blank page.

Every writer knows the empty page, the cursor blinking like a metronome for your guilt. Ask for a first draft of anything and you get one in seconds. It will probably be mediocre. That is the trick working: reacting is easier than creating, and editing is the part humans are good at.

It never tires of explaining.

Ask it to explain anything, at any level, from any angle, at three in the morning, for the ninth time. It does not sigh or glance at the clock. It does not care that you asked the same thing yesterday, and it holds no opinion about it either. Human explainers are wonderful and scarce: they get tired, they charge by the hour, they judge, however gently. A tutor with infinite patience was never affordable before, for anyone. Now the price of admission is typing the question.

A hundred ideas, zero hurt feelings.

Brainstorming with humans is delicate: everyone half-loves their own suggestion, and rejecting one costs social capital. Claude produces twenty ideas, watches you throw away nineteen, and offers twenty more with no attachment to any of them. When ideas cost nothing and nobody sulks, you can be greedy: ask for the safe ones, the strange ones, the terrible-on-purpose ones. Quantity first, taste second. Your taste is the scarce ingredient here, so spend it on choosing, not generating.

YOU:

20 names for my lawnmowing business. puns welcome

CLAUDE:

Lawn and Order. Mow Money
Mow Problems. The Grass
Ceiling. Seventeen more
where those came from.

People who cannot code are coding.

People who cannot program are building small working tools by describing them in plain English. That sentence was science fiction until recently. A spreadsheet cleaner. A page that tracks a hobby. A little script, a saved list of instructions, that renames a thousand files in one go. You explain what you want the way you would explain it to a patient friend, and working code comes back. This book will not teach you programming. Volume Three will show you that, for small things, you no longer need to be taught.

★ Programmers were nervous about this. Then they started doing it too.

The rubber duck talks back.

Programmers keep a strange tradition: when stuck on a bug, explain it out loud to a rubber duck on the desk. The duck says nothing. The explaining does the work, because putting a problem into words forces you to look at it squarely. Claude is the first duck in history that replies. Describe any stuck problem, a budget, a decision you keep circling, and somewhere in the telling you will often solve it yourself. When you do not, the duck asks the one question you had been avoiding.

YOU:

my bread keeps coming out
flat. new yeast, kneaded it
forever, oven checked

CLAUDE:

How cold is your kitchen at
night?

Defuse the angry email.

Somewhere in your drafts folder sits an email you were right not to send. Fury is clarifying to write and expensive to deliver. Paste the furious version in, ask for one that is civil but firm, and you get your argument back with the shrapnel removed. It works in every direction: the stiff note warmed up, the apology that actually sounds like one. This is not deep magic. It is the transformation superpower from the start of this chapter, pointed at the most dangerous text you produce.

YOU:

make this polite but i want him to feel it: "Greg. This is the THIRD reminder."

CLAUDE:

"Hi Greg, following up once more on this." He will feel it.

It shines where iteration is cheap.

Look back across this list and one shape keeps turning up. Drafts, explanations, ideas, small tools, tone: in every case, trying again costs nothing and you are the judge of what is good. The sweet spot has a shape: cheap iteration, your judgment. Hold onto that shape and you can predict where it will delight you before you ask. You can also predict where it will get you into trouble: anywhere a wrong answer is expensive and you cannot tell wrong from right. Which brings us, honestly, to the next chapter.

WHAT IT IS BAD AT

Every superpower ships with known issues.

Sometimes it makes things up.

The superpowers come with a catch, and the biggest is that sometimes Claude states a false thing in a calm, perfectly assured voice. A book that does not exist. A date that is off by a decade. A study nobody ran. The field calls this hallucination: an answer assembled from the shape of true answers rather than from facts. Volume Two explains why it happens. This chapter does something more useful first: it teaches you to expect it and route around it.

YOU:

whats a good book on the history of doorknobs

CLAUDE:

Try "The Turning Point: A History of the Doorknob" by Margaret Hollis.

Confidence and correctness are separate dials.

When a person is unsure, you can hear it: the hedging, the pause, the shrug. Claude's voice has no shrug. A made-up answer arrives in the same steady tone as a solid one, because good answers sound sure. Tone tells you how an answer is dressed, never whether it is true.

Counting letters is weirdly hard.

Ask it to count the letters in a word and it may fumble, which seems absurd for a machine built out of text. Chapter two promised you the reason: it does not read letter by letter, it reads in chunks called tokens, a few characters swallowed whole. Counting letters through tokens is like counting the atoms in a sandwich while eating it: wrong layer entirely. Arithmetic on big numbers fails the same way, estimated rather than computed, the way you estimate a tip. Newer models do better, and some reach for a calculator without saying so. The lesson outlives the fixes: arithmetic that matters gets checked.

YOU:

quick, how many b's in
bubblegum

CLAUDE:

Two.

★ The answer is three. It saw two chunks and zero letters.

It cannot see your life.

Claude does not know your calendar, your inbox, your codebase, or what your manager said in the meeting. Unless you paste those things into the conversation, they do not exist for it. Ask about them anyway and it may guess, smoothly, because guessing smoothly is its native talent. A guess about the world is a hallucination. A guess about your world is a hallucination with your name in it, and those sting worse. The fix costs nothing: paste the thing. Context you can see and it cannot is the biggest source of answers that feel strangely off.

Yesterday has not happened yet.

You met the cutoff in chapter four: training ends on a date, and nothing after that date made it into the machine's head. It bears repeating here, because people rediscover this failure weekly. Ask about last night's game, this morning's price, or the newest version of anything, and an unassisted Claude answers from a world frozen months ago, sometimes without mentioning the freeze. The cure is one sentence long: ask it to search. It then reads live pages and cites them, and the frozen world thaws for exactly as long as you need it to.

It would rather agree with you.

Helpfulness has a shadow. A machine trained to be useful learns to go along: it adopts your framing, finds merit in your plan, softens the bad news. Ask a leading question and you will often get the answer you led it to. This is not deceit, it is politeness at industrial strength. The counter is to ask for the pushback explicitly. Invite the counterargument and it delivers, often sharply, because honesty is in its training too. You just have to signal that today, honesty is the kind of help you want.

YOU:

ok be honest with me. is my
umbrella subscription
startup a bad idea

CLAUDE:

Honestly: the margins leak.
People lose umbrellas, they
do not subscribe to them.

★ Asked as 'isn't this brilliant?', the same plan got a warmer review.

Long conversations drift off course.

Somewhere around the hundredth message, the rules start to bend. The format you set in message three erodes. The voice you asked for slips. The thing you said never to do gets done. As a chat grows, your early instructions sit further and further back in the pile of text the machine rereads on every turn, and distant text pulls less weight. Restate what matters in one line the moment you notice drift, or start a fresh chat with the new chat button and carry over only the good parts. The full mechanics of the pile are Volume Two's whole subject.

★ The pile has a proper name. It is printed on the cover of Volume Two.

None of this is a dealbreaker.

Read this chapter back and notice what it is: not a warning label, a user manual, the kind nobody prints. You now know the machine can bluff, miscount letters, misjudge its own certainty, miss your private world, lag the news, flatter your framing, and forget early promises in long chats. That puts you ahead of most people who use it every day. Knowing where the potholes are does not make you drive less, it makes you drive well. There is one habit that handles all seven potholes at once. It is called verification, and it is the next chapter.

TRUST, BUT VERIFY

The oldest security advice, now in paperback.

Trust in proportion to the stakes.

Constant vigilance is not the answer: nobody double-checks everything, and nobody should. A dinner recipe from Claude needs no verification, because if the paprika is wrong, dinner is slightly worse. A medication dose, a visa deadline, a tax figure: check those every time, no exceptions. Match your checking effort to the price of being wrong. You already run this math with people. You take a stranger's restaurant tip on faith and get a second opinion before surgery. The machine gets the same treatment.

★ If dinner at your house is high stakes, verify the recipe too.

One source is a rumor.

Journalists have a rule older than computers: a fact is not a fact until a second, independent source confirms it. Borrow it. When an answer matters, confirm it somewhere Claude did not write: the official site or the person involved. Boring, ancient, undefeated: what you want from a rule.

Errors hide in the leaps.

Ask Claude to walk you through its reasoning. Not because the steps are guaranteed correct, but because a wrong answer usually survives by skipping the exact step where it went wrong. Laid out one move at a time, the bad move becomes visible, often to you, sometimes to Claude itself. Reading the steps also teaches you the topic. You came to check an answer and left understanding the subject.

YOU:

wait walk me thru how u got that number

CLAUDE:

Step two assumed 30 days in the month. It has 31. The answer changes.

Confidence is a forecast, not a warranty.

You can ask Claude how sure it is. The answer carries real signal: it is noticeably better than nothing at flagging shaky ground. But treat that self-assessment like a weather forecast. A forecast of rain is worth carrying an umbrella for, and it is still sometimes wrong about rain. Remember the lesson from the last chapter: the confident voice and the correct answer are not wired together. Asking how sure it is narrows the gap. Nothing closes it.

Make it argue the other side.

After Claude gives you an answer, ask it to make the case that the answer is wrong. It will, without ego, and often with alarming skill. This is the fastest weak-plank detector available: helpfulness pulls it toward agreeing with its own work, this request pulls the other way, and whatever survives both pulls has earned some trust. Security people call the move red teaming: hiring your own attacker before a real one shows up. You just hired one for free.

YOU:

ok now make the case that
your answer is wrong

CLAUDE:

Gladly. It assumes your
flight lands on time. Your
airline's record disagrees.

A clickable citation beats a vibe.

When a claim matters, ask where it comes from. Answers from memory arrive with no receipts, because memory has none: it is reconstruction, not retrieval. But you can ask Claude to search the web and cite what it finds, and then every claim arrives with a link you can open and judge for yourself. This is the difference between the two modes from chapter four, memory and research, now used deliberately. Freshness was one reason to ask for research. Receipts are the other.

YOU:

wheres that stat from

CLAUDE:

Memory, which is not a source. Want me to search for a real one?

You are the editor in chief.

Every newsroom runs on a division of labor: reporters draft fast and file plenty, and an editor decides what is true enough and good enough to print. That editor is you now. Claude drafts, explains its steps on request, argues both sides on command. You sign the page, and the signature never transfers. The arrangement does not expire as you improve: in Volume Three you will run tools and agents, programs that act on your say-so, and you will still be the one who signs. With the checking habit installed, you can hand the machine real work with a clear conscience, and the next chapter is nothing but recipes.

RECIPES

Steal these. That is what they are for.

Paste the draft. Say the goal.

Paste the email you have been rewriting for twenty minutes. Say who the person is to you and what you want to happen: my gym, I want the membership canceled this time, polite but immovable. Then ask for the version most likely to get that reply. The recipe ends there. The draft was never your problem. Your problem was that you knew the goal and the email did not, and now the two have finally been introduced.

Explain it, then quiz me.

One line turns an answer into a lesson: explain this like I am new to it, then quiz me with three questions. The explanation is the easy half. The quiz is where the learning happens, because being asked forces you to reach for the idea instead of nodding along with it. Get a question wrong and it explains the exact thing you missed, then asks again. You have built a loop: teach, test, patch the gap, repeat until the gap is gone.

YOU:

explain how vaccines work
like im new, then quiz me
with 3 questions

CLAUDE:

A wanted poster for the
virus. Question one: why
does the second exposure go
better?

Constraints in, plan out, argue after.

Trips, parties, budgets, the week ahead: anything with moving parts. Hand over the constraints, meaning the dates, the money, who is coming, and the one thing you refuse to do. A full plan comes back in seconds, and parts of it will be wrong for you. Push back on everything you hate: too packed, no museums, we are not morning people. Each objection reshapes the plan, and somewhere around the third round you notice you are not filling in a form, you are having a conversation about what you want.

It also translates the dialects of life.

Between languages first: it reads and writes most of the big ones and a good number of the small ones. But the second kind of translation is the one you will use weekly. Legalese into English. Doctor into human. The letter from the tax office into what they actually want from you. Every profession speaks a dialect built for precision, and to outsiders the same dialect reads as a locked door. Paste the paragraph and ask what it says in plain words, and whether you need to do anything about it.

YOU:

what does "the party of the first part hereby indemnifies" even mean

CLAUDE:

If things go wrong, you pay. That clause deserves a lawyer's eyes before you sign.

The quiz comes from your notes.

The explainer teaches you things you never knew. The study buddy drills the things you are supposed to know by Tuesday. Paste your notes, your flashcards, the chapter summary, then say: quiz me on this, get harder when I am right, explain whatever I miss. It builds the questions from your material, not its own, so you are rehearsing the exam you will sit. Wrong answers turn into small lessons on the spot, and the difficulty climbs as you improve, which is exactly what a good tutor does and a stack of flashcards never will.

It will not decide for you. Good.

Lay the choice out honestly: take the job or stay, buy or rent. Ask for the strongest case for each side, argued like it means it. Then the follow-up: what am I not considering? It never picks for you. It does not have to live with the outcome. You do.

Outline together, draft alone, edit together.

Writers worry that using the machine means sounding like the machine, and if you let it draft everything, they are right. So split the work. Outline together: structure is where it is fast and you are slow. Draft alone: the sentences come from you, in your voice, mess and all. Then edit together: it tightens, spots the sagging middle, questions the paragraph you secretly knew was weak. The words stay yours from first to last. You have borrowed speed at both ends.

YOU:

heres my draft. dont rewrite it, just tell me whats weak

CLAUDE:

Paragraph three buries your best line. Move it up and cut the apology in front of it.

The infinite was never the machine.

Seven recipes, and none of them were the point. The point is that you can now invent your own, because a recipe is just a conversation somebody bothered to write down, and your life is full of conversations nobody has had yet. That is the infinite this volume was named for. It lives in the unasked. One catch, though: while you talk, the machine is quietly forgetting. How, why, and what to do about it fills the next book. Volume Two: THE CONTEXT WINDOW.

NEXT

**But there is a catch.
The machine forgets.**

Volume Two: THE CONTEXT WINDOW

Who wrote this?

Suraaj Chandra Jha is a hacker in the old, joyful sense of the word: a person who delights in understanding systems and making them do wonderful things. He has been explaining computers to humans, and humans to computers, for longer than he will admit. This trilogy is free because the best tricks have always been passed around for free.

ONE LAST THING

Pass it on.

If this book taught you one thing, send it to one person. That is the license fee. All three volumes live at davidjairaj.github.io/sharing-for-a-friend, free, always.